

10th Workshop in Decisions, Games and Logic: Ethics, Statistics, and Fair AI

June 10-12, 2019, California Institute of Technology.

All talks will be held in Baxter Room B125.

Participation is free and open to all. If you would like to attend and are not on the program please RSVP to caltech.dgl2019@gmail.com by June 1, 2019.

Directions to Caltech: <https://www.caltech.edu/about/visit/directions>

Campus parking (\$6/day): <http://parking.caltech.edu/Parking>

Schedule (DRAFT)

Monday, June 10

9:00-9:45 Fairness Without Harm: Decoupled Classifiers with Preference Guarantees
Berk Ustun (Harvard University), Yang Liu (UC Santa Cruz), David Parkes (Harvard University)

9:45-10:30 Learning Optimal Fair Policies
Razieh Nabi (Johns Hopkins), Daniel Malinsky (Johns Hopkins), Ilya Shpitser (Johns Hopkins)

10:30-11:00 Break

11:00-12:15 Title TBD
Ilya Shpitser (Johns Hopkins)

12:15-1:30 Lunch

1:30-2:45 Measuring Algorithmic Fairness
Deborah Hellman (University of Virginia)

2:45-3:30 Algorithmic Injustice Beyond Discriminatory Harm
Annette Zimmermann (Princeton University)

3:30-4:00 Break

4:00-4:45 Counterfactuals as Causal Fictions
Lily Hu (Harvard University)

4:45-6:00 Can we Trust a Biased AI?
David Danks (Carnegie Mellon)

6:30-8:30 Workshop Dinner (Location TBD)

Tuesday, June 11

9:00-9:45 Preference Informed Fairness

Aleksandra Korolova (USC), Michael Kim (Stanford), Gal Yona (Weizmann Institute), Guy Rothblum (Weizmann Institute)

9:45-10:30 The Aspirations and Shortcomings of Causal Approaches to Algorithmic Fairness

Sina Fazelpour (University of British Columbia) and Zachary C. Lipton (Carnegie Mellon)

10:30-11:00 Break

11:00-12:15 The Measure and Mismeasure of Fairness: A Critical Review of Fair Machine Learning

Sharad Goel (Stanford University)

12:15-1:30 Lunch

1:30-2:45 Moral Responsibility, Blameworthiness and Intention: In Search of Formal Definitions

Joseph Halpern (Cornell University)

2:45-3:30 Principles for Public Sector Data Analytics

Daniel Weissglass (CUNY)

3:30-4:00 Break

4:00-5:15 Towards Practical Notions of Individual Fairness

Aaron Roth (University of Pennsylvania)

5:15-6:00 Tracking and Improving Information in the Service of Fairness

Michael Kim (Stanford University)

6:30-8:30 Workshop Dinner (Location TBD)

Wednesday, June 12

9:00-9:45 A Game Semantics for Tolerance

Ali Abasnezhad (LMU Munich)

9:45-10:30 Fairness in the Context of Algorithmic Decision Making: Towards a Map of Ethical Concerns

Florian Ostmann (Turing Institute)

10:30-11:00 Break

11:00-11:45 How AI Might Help us Die With Dignity and Fairness

Russell DiSilvestro (Cal State Sacramento), Daniel Weijers (Univ. of Waikato), and Nicholas Williams (Cornerstone OnDemand, Inc.)

11:45-12:30 Analysis of Maximally Approved Police Forces

Jamie Vinson (Caltech)